

Image Generation using Stable Diffusion & Comfy UI

A Project Report

submitted in partial fulfillment of the requirements

of

AICTE Internship on AI: Transformative Learning
with

TechSaksham – A joint CSR initiative of Microsoft & SAP

by

Umang Aggarwal,

umangddu5@gmail.com

Under the Guidance of

Jay Rathod

ACKNOWLEDGEMENT

I would like to express my heartfelt gratitude to everyone who contributed to the successful completion of this project. Their unwavering support, guidance, and encouragement have been invaluable throughout this journey.

First and foremost, I extend my sincere thanks to my mentor, **Jay Rathod**, for his insightful guidance, constructive feedback, and constant encouragement. His expertise has played a crucial role in shaping the direction of this project. I also deeply appreciate the support of **Adarsh P**, whose valuable inputs have helped refine my approach and enhance my understanding of AI-driven image generation.

I would also like to express my gratitude to **Pavan Kumar Sir** for his continuous support throughout the internship. His assistance in administrative tasks, including communications and offer letters, ensured a smooth and hassle-free learning experience.

Furthermore, I extend my appreciation to **AICTE** and the **TechSaksham initiative by Microsoft & SAP** for providing me with this incredible opportunity to deepen my knowledge of AI technologies. This internship has been instrumental in broadening my understanding of generative AI and its real-world applications.

Lastly, I acknowledge the contributions of the **open-source community** and the researchers behind **Stable Diffusion and ComfyUI**. Their innovative work and publicly available resources have laid the foundation for this project, enabling exploration and implementation of AI-powered image generation.

This project has been a significant learning experience, and I am immensely grateful to everyone who has supported me throughout this journey.

ABSTRACT

Recent advancements in generative AI have significantly improved image synthesis capabilities, with Stable Diffusion emerging as a leading model for high-quality text-to-image generation. However, utilizing Stable Diffusion effectively requires intricate configurations and parameter tuning, making it challenging for users without technical expertise. This project integrates Stable Diffusion with ComfyUI, a node-based interface that simplifies workflow customization and enhances usability.

The primary focus of this work is to optimize image generation by leveraging ComfyUI's visual approach, allowing for better control over the synthesis process. The methodology involves setting up the Stable Diffusion model, fine-tuning generation parameters, and experimenting with diverse configurations to achieve high-fidelity results. This integration bridges the gap between technical complexity and user-friendly design, enabling a broader audience to harness the power of AI-driven image generation effectively.

In the modern digital ecosystem, online collaboration and communication play a vital role in knowledge sharing and problem-solving. This project presents a chat room feature that enables users to form discussion groups and engage in real-time conversations. A distinguishing aspect of this system is the integration of the Gemini AI Model 1.5, which enhances user interactions by providing AI-driven insights. Users can seamlessly request AI assistance by incorporating the '@ai' keyword within their messages, ensuring relevant and structured responses to their queries.

The chat room feature is designed with user-friendliness in mind, incorporating clear project setup guidelines to facilitate smooth implementation. By leveraging advanced AI capabilities, this system offers an intuitive platform where users can engage in meaningful discussions, seek solutions, and collaborate efficiently. The AI model assists in refining prompts and generating context-aware responses, significantly improving the overall communication experience.

Furthermore, the project ensures that AI interactions are seamlessly integrated into the chat environment without disrupting the natural flow of discussions. The combination of real-time messaging and AI-driven insights fosters a dynamic and interactive communication space, enhancing user engagement and problem resolution. This implementation not only streamlines online discussions but also demonstrates the potential of AI-powered assistance in collaborative platforms.

TABLE OF CONTENT

Abstract		I
Chapter 1. Introduction		1
1.1 Problem Statement		1
1.2 Motivation		1
1.3 Objectives		2
1.4 Scope of the Project		2
Chapter 2. Literature Survey		4
2.1 Literature Review		4
2.2 Existing Models and Techniques		4
2.3 Gaps in Existing Solutions		4
Chapter 3. Proposed Methodology		6
3.1 System Design		6
3.2 Requirement Specification		7
Chapter 4. Implementation and Results		9
4.1 Snapshots of result		9
4.2 GitHub link for code		10
Chapter 5. Discussion and Conclusion		11
5.1 Future work		11
5.2 Conclusion		11
References		13

LIST OF FIGURES

Figure No.	Figure Caption	Page No.
Figure 1	Mystical Underwater City	9
Figure 2	Time Traveler's Ship in the Sky	9
Figure 3	Giant Tree with Floating Lanterns	9
Figure 4	Futuristic Market on a Distant Planet	9
Figure 5	Floating City Above the Clouds	10
Figure 6	Ice Dragon in a Snowy Mountain Range	10
Figure 7	Human-AI Hybrid in a Futuristic City	10
Figure 8	AI Avatar in Virtual Reality	10

CHAPTER 1

Introduction

1.1 Problem Statement:

The demand for AI-generated images is growing across various fields, including design, marketing, and entertainment. Stable Diffusion, a powerful text-to-image model, offers high-quality image synthesis but requires complex configurations, parameter tuning, and workflow optimization. This makes it challenging for users without deep technical expertise to fully leverage its potential.

Stable Diffusion provides impressive image generation capabilities, but its complexity poses challenges for users lacking technical expertise. Current implementations require detailed understanding of model parameters, computational requirements, and configuration settings, limiting accessibility. Additionally, existing interfaces may not provide sufficient control over the generation process, leading to suboptimal results.

This project aims to simplify the utilization of Stable Diffusion by integrating it with ComfyUI, a visual workflow tool that enhances customization and ease of use. The goal is to provide an intuitive interface that allows users to fine-tune parameters efficiently while maintaining high-quality outputs. By addressing these usability issues, the project seeks to bridge the gap between advanced AI models and practical, user-friendly applications.

1.2 Motivation:

The rapid evolution of generative AI technologies has significantly transformed how we approach content creation, especially in areas like digital art, media production, and design. However, despite the potential of tools like Stable Diffusion, many users face challenges in utilizing these technologies due to their complexity and technical nature. This project is driven by the need to simplify access to AI-driven image generation, making it more user-friendly and accessible to a broader audience, even those with limited technical experience.

The main motivation behind this project is to create a seamless and intuitive experience for users by providing an interface that eliminates the steep learning curve typically associated with AI models. Tools like ComfyUI offer a node-based structure, which simplifies the task of configuring and generating images, allowing users to focus on creativity rather than dealing with intricate model parameters.

This project is also motivated by the desire to explore how AI can be applied to various domains in a more practical, accessible manner:

- **Creative Industries:** Artists and designers can use this technology to produce unique visuals quickly, enhancing productivity and creativity.
- **Marketing & Branding:** Companies can create customized promotional content more efficiently, helping them stay ahead in competitive markets.

- **Gaming:** Game developers can use AI-generated assets to streamline the creation of concept art and game environments.
- **Educational Tools:** AI-powered design can serve as an educational resource for students and researchers interested in learning about AI and its creative applications.
- **Content Creation:** Bloggers, influencers, and content creators can use AI to generate fresh visuals for their social media platforms or websites.

By offering an optimized, simplified solution for interacting with generative AI, this project aims to unlock the full creative potential of users, encouraging them to explore and experiment with AI technologies without the need for advanced technical skills. The goal is to make AI-driven content creation accessible to all, fostering innovation and creativity across industries.

1.3 Objective:

The primary objective of this project is to develop an intuitive and powerful system for generating high-quality AI images from textual descriptions. By integrating **Stable Diffusion** with a **ComfyUI** interface, the goal is to simplify the image generation process and enhance user experience.

Key objectives of the project include:

1. **Stable Diffusion Integration:** Implementing Stable Diffusion to generate high-quality images from textual prompts.
2. **ComfyUI Integration:** Utilizing ComfyUI to provide a visual, node-based interface that simplifies workflow customization and enhances ease of use.
3. **Parameter Optimization:** Optimizing generation parameters, including denoising steps, guidance scale, and sampling methods, to improve image output quality.
4. **User Control:** Ensuring users can refine and modify outputs effectively, providing better control over image generation.
5. **Semantic Accuracy Evaluation:** Using CLIP to evaluate generated images against input prompts, improving semantic alignment and accuracy.
6. **AI Imagery Applications:** Exploring potential applications of AI-generated imagery in fields such as digital art, content creation, and game design.

By meeting these objectives, the project aims to make generative AI tools more accessible, foster creativity, and promote innovative applications across various industries.

1.4 Scope of the Project:

The scope of this project is centered around leveraging cutting-edge AI technologies to transform textual input into high-quality images. By combining **Stable Diffusion** with a visual

interface provided by **ComfyUI**, this project aims to make AI-powered image generation accessible and customizable for a wide range of users.

Key aspects of the project scope include:

1. **Text-to-Image Creation:** The project will allow users to generate realistic images based on written descriptions. This can be useful for artists, content creators, and marketers.
2. **Easy-to-Use Interface:** With **ComfyUI**, users can interact with a simple, visual interface that makes it easy to control and adjust how the images are generated, without needing technical skills.
3. **Customization Options:** The project will offer settings to control important aspects of image creation, like denoising and sampling, so users can fine-tune the results to their liking.
4. **Ensuring Accuracy:** By using **CLIP**, the project will check if the generated images match the input text, improving the accuracy and relevance of the results.
5. **Real-World Applications:** The project also aims to explore how AI-generated images can be used in different fields like gaming, advertising, and digital art.

Overall, the project will make it easier for users to create customized images using AI, with a focus on creativity and user control.

1.4.1 Limitations:

While this project offers an innovative way to generate images using AI, there are some challenges and limitations to consider:

- **Hardware Dependencies** – Running Stable Diffusion efficiently requires high computational power (e.g., a GPU with sufficient VRAM). Performance may be limited on lower-end systems.
- **Limited Realism in Generated Images** – While Stable Diffusion produces high-quality images, it may struggle with highly detailed or complex scenes without fine-tuning.
- **Text-to-Image Ambiguity** – The model's interpretation of prompts may not always align perfectly with user expectations, requiring multiple iterations.
- **No Real-time Generation** – The process involves computationally intensive steps, making real-time image generation impractical for some use cases.
- **Dependency on Pre-trained Models** – The project relies on existing Stable Diffusion models, and any significant improvements require additional fine-tuning or external datasets.

Despite these limitations, the project provides a practical and efficient approach to AI-based image generation, offering a balance between usability, control, and creative flexibility.

CHAPTER 2

Literature Survey

2.1 Literature Review

AI-driven image generation has advanced significantly, especially with models like **Stable Diffusion**. Stable Diffusion uses **diffusion models** to generate high-quality images from text prompts, offering better control and faster results compared to older models like **GANs (Generative Adversarial Networks)** and **VAEs (Variational Autoencoders)**.

Text-to-Image Models such as **CLIP**(Contrastive Language-Image Pretraining) help assess the alignment between generated images and input descriptions, though challenges like prompt ambiguity still exist. **Stable Diffusion** improves upon this by providing more accurate and customizable outputs.

The project utilizes **ComfyUI**, a visual interface for Stable Diffusion, allowing easy parameter adjustments without coding expertise. This makes AI image generation accessible to a broader audience, especially in fields like **digital art**, **content creation**, and **game design**.

2.2 Existing Models and Techniques

- **GANs (Generative Adversarial Networks)** – Earlier generative models like StyleGAN and BigGAN produced high-quality images but often suffered from mode collapse and instability in training.
- **VAEs (Variational Autoencoders)** – Used for latent space-based image generation, but the outputs tend to be blurry compared to diffusion models.
- **DALL·E (OpenAI)** – A transformer-based model that generates images from text but is not as open-source or flexible as Stable Diffusion.
- **CLIP (Contrastive Language-Image Pretraining)** – Developed by OpenAI, CLIP is used to evaluate and refine generated images based on their alignment with text prompts.

2.3 Gaps in Existing Solutions:

- **Complexity in Implementation** – Many AI-based image generation tools require technical expertise, such as knowledge of programming or command-line execution, and often involve complicated parameter adjustments. This project addresses this gap by introducing **ComfyUI**, a user-friendly, visual interface that makes the process simpler and more accessible to a wider range of users.
- **Limited Control Over Outputs** – While diffusion models like Stable Diffusion are capable of generating impressive images, achieving specific styles or detailed attributes can be a challenge. This project improves user control by fine-tuning generation parameters, and


utilizing **CLIP-based evaluations** to ensure more accurate and consistent outputs that align better with textual descriptions.

- **Computational Limitations** – High-quality AI image generation typically requires powerful hardware, especially GPUs, which can be a barrier for many users. This project aims to overcome this limitation by exploring optimization techniques and efficient configurations that enable users with moderate hardware setups to also benefit from the technology.
- **Text-to-Image Ambiguity** – Sometimes, AI models like Stable Diffusion may misinterpret the provided prompts, resulting in undesired or inaccurate images. This project focuses on refining workflows and experimenting with prompt engineering strategies to improve the clarity and accuracy of text-to-image translations, enhancing the overall consistency of the generated images.

CHAPTER 3

Proposed Methodology

3.1 System Design



This diagram represents the Stable Diffusion-based image generation pipeline using ComfyUI. It illustrates the sequence of steps involved in generating an AI-generated image from a text prompt using a node-based approach.

Below is a detailed breakdown of each node and its role in the process:

- **Load Checkpoint (Model Selection):** This node loads the Stable Diffusion model checkpoint (v1-5-pruned-emaonly-fp16), which is used for image generation.

It also loads the CLIP model for text encoding and the VAE (Variational Autoencoder) for decoding the generated latent image.

- **CLIP Text Encode (Prompt):**

Two CLIP Text Encode nodes are used:

- **Positive Prompt:** "A futuristic city at sunset" – This describes the desired output.

- **Negative Prompt:** "text, watermark" – Specifies unwanted elements in the generated image.

These nodes convert the textual description into a latent space representation, guiding the model's image generation process.

- **Empty Latent Image (Image Size Definition):** This node initializes a blank latent space with dimensions 512x512 pixels and a batch size of 1. It provides the canvas where the generated image will take shape.
- **KSampler (Image Generation Process):** This node is responsible for sampling and refining the latent image using the Stable Diffusion model.

Key parameters:

- Seed: 975502189541890 (ensures reproducibility if reused).
- Steps: 20 (controls the number of denoising iterations).
- CFG (Classifier-Free Guidance Scale): 8.0 (adjusts prompt adherence).
- Sampler: Euler (determines how noise is removed at each step).
- Scheduler: Normal (sets the denoising schedule).
- Denoise: 1.00 (controls the level of noise removal).
- **VAE Decode (Latent to Image Conversion):** After the latent image is generated, this node decodes it into a visible image using the Variational Autoencoder (VAE). This step translates the compressed latent representation into a high-resolution image.
- **Save Image (Output Storage):** The final image is saved with the filename prefix "ComfyUI". This allows users to store and view the AI-generated image.

3.2 Requirement Specification

To implement the Stable Diffusion Image Generation project using ComfyUI, we require a combination of hardware and software components for optimal performance.

3.2.1 Hardware Requirements:

- **Processor:** Minimum Intel i5 / Ryzen 5 (Recommended: Intel i7 / Ryzen 7 or higher)

- **GPU:**
 - Can run in CPU mode but will be significantly slower
 - NVIDIA RTX 3060 (Minimum), RTX 3080 / 4090 (Recommended) for faster image generation
- **RAM:** At least 16GB RAM (Recommended: 32GB or more)
- **Storage:** SSD (Minimum 512GB) for quick model loading and caching
- **Power Supply:** Adequate wattage to support high-end GPUs

3.2.2 Software Requirements:

- **Operating System:** Windows 10/11 or Linux (Ubuntu preferred) or macOS (M1/M2 with workarounds)
- **Stable Diffusion Model:** Stable Diffusion v1.5, SDXL 1.0, or fine-tuned models
- **ComfyUI:** Node-based workflow for image generation
- **Python:** Python 3.10+ for scripting and automation
- **Version Control:** Git and GitHub for managing code and workflow

CHAPTER 4

Implementation and Result

4.1 Snap Shots of Result:




Figure 1: Mystical Underwater City



Figure 2: Time Traveler's Ship in the Sky



Figure 3: Giant Tree with Floating Lanterns



Figure 4: Futuristic Market on a Distant Planet




Figure5: Floating City Above the Clouds



Figure5: Ice Dragon in a Snowy Mountain Range




Figure5: Human-AI Hybrid in a Futuristic City



Figure5: AI Avatar in Virtual Reality

4.2 GitHub Link for Code:

<https://github.com/umangaggarwal55/Image-Generation-using-stable-diffusion-Comfy-UI.git>

CHAPTER 5

Discussion and Conclusion

5.1 Future Work:

This project successfully integrates **Stable Diffusion** with a **ComfyUI** interface to simplify the process of AI-based image generation, providing an accessible tool for both beginners and experts. By optimizing key parameters and using **CLIP** evaluations, it improves the alignment between generated images and text prompts, offering greater control over output quality.

While the current implementation of Stable Diffusion with ComfyUI demonstrates strong image generation capabilities, there are several areas for improvement and future development:

1. **Model Efficiency & Optimization:**
 - Develop optimized, lighter models that require less computational power without compromising image quality.
 - Explore diffusion techniques like **LoRA (Low-Rank Adaptation)** to enable fine-tuning with fewer resources.
2. **User Interface & Workflow Enhancements:**
 - Create a more user-friendly web-based UI to facilitate easier customization and streamline workflow automation.
 - Introduce real-time preview adjustments, allowing users to modify prompts and settings dynamically during image generation.
3. **Expanding Use Cases:**
 - Investigate extending Stable Diffusion's functionality to video generation, incorporating frame interpolation techniques for fluid transitions.
 - Enable multi-modal input capabilities, such as combining text, sketches, or rough images for more nuanced prompts.
4. **Addressing GPU Dependency & Performance on CPU:**
 - Currently, Stable Diffusion and ComfyUI rely on high-performance NVIDIA GPUs for optimal performance due to CUDA and Tensor core acceleration. When running without a GPU, the model significantly slows down.
 - **Performance Comparison:**
 - With **NVIDIA GPU (RTX 3060 or higher)**: Image generation takes just a few seconds.

- Without **GPU (CPU mode)**: Generation can take several minutes per image, highlighting the need for optimization in CPU mode to improve efficiency.

These areas of focus will enhance the performance, flexibility, and accessibility of the project, expanding its potential use cases and enabling wider adoption.

5.2 Conclusion:

This project successfully demonstrates the power of AI in generating high-quality images using Stable Diffusion combined with ComfyUI. By utilizing Stable Diffusion, a robust model for creating images from text, and integrating it with ComfyUI, the project offers an easy-to-use, customizable method for generating impressive images.

The implementation highlights the role of deep learning techniques, such as UNet-based models and CLIP-guided diffusion, in producing diverse and appealing visuals. It also explores methods like CFG (Classifier-Free Guidance) and different scheduler variations to enhance the quality and consistency of generated images.

However, the project also reveals some challenges, such as the high computational requirements for optimal performance, especially the need for GPU acceleration. Running the model on a CPU leads to slower processing times, making it less accessible for users without powerful hardware.

In conclusion, this project adds value to the field of AI-generated images by showcasing the effectiveness of diffusion models. With future improvements like model optimization, hardware adjustments, and better control features, this work can contribute to more efficient and creative solutions in AI image generation.

REFERENCES

- [1]. **J. Ho, A. Jain, and P. Abbeel**, "Denoising Diffusion Implicit Models," *Proceedings of NeurIPS*, 2020. [Online]. Available: <https://arxiv.org/abs/2006.11239>
- [2]. **M. Zhang, Y. Yu, J. S. Choi, and L. Yu**, "Latent Diffusion Models for Text-to-Image Generation," *ACM Transactions on Graphics (TOG)*, vol. 41, no. 5, 2022. [Online]. Available: <https://dl.acm.org/doi/10.1145/3481618>
- [3]. **R. Rombach**, "Understanding the Stable Diffusion Model and Its Applications," *Medium*, Sep. 2022. [Online]. Available: https://medium.com/@rombach_r/stable-diffusion-explained
- [4]. **R. Narayan**, "The Future of Text-to-Image Models: A Study of CLIP and Stable Diffusion," *AI Review Journal*, 2024. [Online]. Available: <https://ai-review.com/articles/clip-stable-diffusion>
- [5]. **A. Sharma, M. Singh, and R. Agrawal**, "Evaluating CLIP for Better Image-Text Alignment in Stable Diffusion," *arXiv preprint arXiv:2311.04109*, 2023. [Online]. Available: <https://arxiv.org/abs/2311.04109>
- [6]. **Hugging Face**, "Stable Diffusion: Next Generation Image Synthesis," *Hugging Face Documentation*, 2022. [Online]. Available: <https://huggingface.co/stabilityai/stable-diffusion>